

Peeking into Pre-Defined Cluster Roles

In this lesson, we will look into all the pre-defined cluster roles.

WE'LL COVER THE FOLLOWING ^

- Switching from John to Us
- Looking into Roles and Cluster Roles
 - Looking into `view`
 - Looking into `edit`
 - Looking into `admin`
 - Looking into `cluster-admin`

Switching from John to Us

John is frustrated. He can access the cluster, but he is not permitted to perform any operation. He cannot even list the Pods. Naturally, he asked us to be more generous and allow him to “play” with our cluster.

Since we are not taking anything for granted, we decided that the first action should be to verify John’s claim. Is it true that he cannot even retrieve the Pods running inside the cluster?

Before we move further, we’ll stop impersonating John and go back to using the cluster with god-like administrative privileges granted to the `minikube` user.

```
kubectl config use-context minikube
kubectl get all
```



Now that we switched to the `minikube` context (and the `minikube` user), we regained full permissions, and `kubectl get all` returned all the objects from the `default` Namespace.

Let's verify that John indeed cannot list Pods in the `default` Namespace.

We could configure the same certificates as those he's using, but that would complicate the process. Instead, we'll use a `kubectl` command that will allow us to check whether we could perform an action if we would be a specific user.

```
kubectl auth can-i get pods --as jdoe
```



The response is `no`, indicating that `jdoe` cannot `get pods`. The `--as` argument is a global option that can be applied to any command. The `kubectl auth can-i` is a “special” command. It does not perform any action but only validates whether an operation could be performed. Without the `--as` argument, it would verify whether the current user (in this case `minikube`) could do something.

Looking into Roles and Cluster Roles

We already discussed Roles and ClusterRoles briefly. Let's see whether any are already configured in the cluster or the `default` namespace.

```
kubectl get roles
```



The output reveals that `no resources` were `found`. We do not have any Roles in the `default` Namespace. That was the expected outcome since a Kubernetes cluster comes with no pre-defined Roles. We'd need to create those we need ourselves.

How about Cluster Roles? Let's check them out.

```
kubectl get clusterroles
```



This time we got quite a few resources. Our cluster already has some Cluster Roles defined by default. Those prefixed with `system:` are Cluster Roles reserved for Kubernetes system use. Modifications to those roles can result in non-functional clusters, so we should not update them. Instead, we'll skip system Roles and focus on those that should be assigned to users.

The output, limited to Cluster Roles that are meant to be bound to users, is as follows (you can get the same result through `kubectl get clusterroles | grep -v system`).

NAME	AGE
admin	1h
cluster-admin	1h
edit	1h
view	1h



Looking into `view`

The Cluster Role with the least permissions is `view`. Let's take a closer look at it.

```
kubectl describe clusterrole view
```



The **output**, limited to the first few rows, is as follows.

```
Name:      view
Labels:    kubernetes.io/bootstrapping=rbac-defaults
Annotations: rbac.authorization.kubernetes.io/autoupdate=true
PolicyRule:
  Resources            Non-Resource URLs  Resource Names  Verbs
  -----
  bindings             []                 []              [get list watch]
  configmaps           []                 []              [get list watch]
  cronjobs.batch       []                 []              [get list watch]
  daemonsets.extensions []                 []              [get list watch]
  deployments.apps     []                 []              [get list watch]
  ...
```



It contains a long list of resources, all of them with the `get`, `list`, and `watch` verbs. It looks like it would allow users bound to it to retrieve all the resources. We have yet to validate whether the list of resources is truly complete.

For now, it looks like an excellent candidate to assign to users that should have very limited permissions. Unlike Roles that are tied to a specific Namespace, Cluster Roles are available across the whole cluster. That is a significant difference that we'll exploit later on.

Looking into `edit`

Let's explore another pre-defined Cluster Role

Let's explore another pre-defined Cluster Role.

```
kubectl describe clusterrole edit
```



The **output**, limited to Pods, is as follows.

```
...
pods                [] [] [create delete deletecollection get list patch update watch]
pods/attach         [] [] [create delete deletecollection get list patch update watch]
pods/exec           [] [] [create delete deletecollection get list patch update watch]
pods/log            [] [] [get list watch]
pods/portforward    [] [] [create delete deletecollection get list patch update watch]
pods/proxy          [] [] [create delete deletecollection get list patch update watch]
pods/status         [] [] [get list watch]
...
```



As we can see, the **edit** Cluster Role allows us to perform any action on Pods. If we go through the whole list, we'd see that the **edit** role allows us to execute almost any operation on any Kubernetes object.

It seems like it gives us unlimited permissions. However, there are a few resources that are not listed. We can observe those differences through the Cluster Role **admin**.

Looking into **admin**

```
kubectl describe clusterrole admin
```



If you pay close attention, you'll notice that the Cluster Role **admin** has a few additional entries.

The output, limited to the records not present in the Cluster Role **edit**, is as follows.

```
...
localsubjectaccessreviews.authorization.k8s.io [] [] [create]
rolebindings.rbac.authorization.k8s.io        [] [] [create delete deletecollection get list]
roles.rbac.authorization.k8s.io                [] [] [create delete deletecollection get list]
...
```



The main difference between **edit** and **admin** is that the latter allows us to manipulate Roles and RoleBindings. While **edit** permits us to do

almost any operation related to Kubernetes objects like Pods and Deployments, `admin` goes a bit further and provides an additional capability that allows us to define permissions for other users by modifying existing or creating new Roles and Role Bindings.

The major restriction of the `admin` role is that it cannot alter the Namespace itself, nor it can update Resource Quotas (we haven't explored them yet).

Looking into `cluster-admin`

There is only one more pre-defined non-system Cluster Role left.

```
kubectl describe clusterrole \
  cluster-admin
```

The **output** is as follows.

```
Name:      cluster-admin
Labels:    kubernetes.io/bootstrapping=rbac-defaults
Annotations: rbac.authorization.kubernetes.io/autoupdate=true
PolicyRule:
  Resources Non-Resource URLs Resource Names Verbs
  -----
  *.*      [*]          []          [*]
  *.*      []          []          [*]
```

The Cluster Role `cluster-admin` holds nothing back. An asterisk (`*`) means everything. It provides god-like powers. A user bound to this role can do anything, without any restrictions. The `cluster-admin` role is the one bound to the `minikube` user. We can confirm that easily by executing

```
kubectl auth can-i "*" "*"
```

The output is `yes`. Even though we did not really confirm that the `cluster-admin` role is bound to `minikube`, we did verify that it can do anything.

In the next lesson, we will learn how to create role bindings.

