

Analysis of breadth-first search

How long does breadth-first search take for a graph with vertex set V and edge set E ? The answer is $O(V+E)$ time.

Let's see what $O(V+E)$ time means. Assume for the moment that $|E| \geq |V|$, which is the case for most graphs, especially those for which we run breadth-first search. Then $|V| + |E| \leq |E| + |E| = 2 \cdot |E|$. Because we ignore constant factors in asymptotic notation, we see that when $|E| \geq |V|$, $O(V+E)$ really means $O(E)$. If, however, we have $|E| < |V|$, then $|V| + |E| \leq |V| + |V| = 2 \cdot |V|$, and so $O(V+E)$ really means $O(V)$. We can put both cases together by saying that $O(V+E)$ really means $O(\max(V,E))$. In general, if we have parameters x and y , then $O(x + y)$ really means $O(\max(x, y))$.

(Note, by the way, that a graph is **connected** if there is a path from every vertex to all other vertices. The minimum number of edges that a graph can have and still be connected is $|V| - 1$. A graph in which $|E| = |V| - 1$ is called a **free tree**.)

How is it that breadth-first search runs in $O(V+E)$ time? It takes $O(V)$ time to initialize the distance and predecessor for each vertex ($\Theta(V)$ time, actually). Each vertex is visited at most one time, because only the first time that it is reached is its distance *null*, and so each vertex is enqueued at most one time. Since we examine the edges incident on a vertex only when we visit from it, each edge is examined at most twice, once for each of the vertices it's incident on. Thus, breadth-first search spends $O(V+E)$ time visiting vertices.