

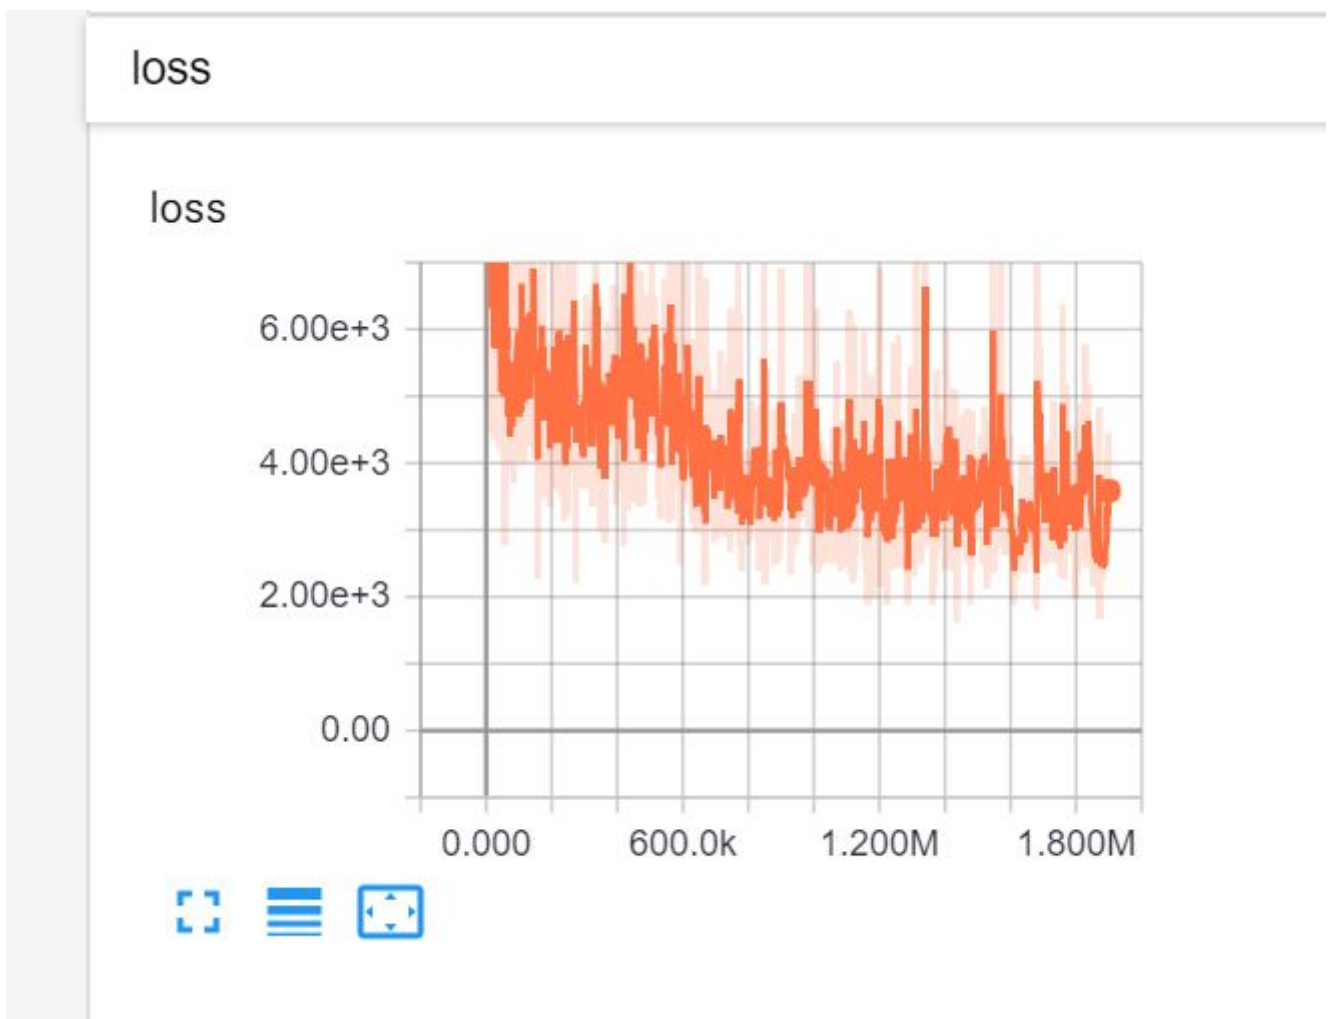
Model Evaluation

Chapter Goals:

- Evaluate the regression model

A. Evaluating with the **Estimator**

We train the model long enough that the loss begins to show signs of convergence (for our 2 hidden layer MLP model, this is around 2M training steps). To get a better look at how the loss is progressing, we can take a look at the TensorBoard training visualizations (see the Model Execution section of [Deep Learning for Industry](#) for more details on TensorBoard).



The TensorBoard line plot showing the loss progression.

We evaluate with the `Estimator` in almost the exact same way we train, with the main difference being that we use the `evaluate` function rather than the `train` function.

The evaluation dataset is contained in the `eval.tfrecords` file (created in the Data Processing Lab). The batch size argument for creating the evaluation TFRecords dataset only affects evaluation speed. A larger batch size can provide a speedup in evaluation, although you have to make sure the batch is small enough to be contained in memory.

We'll use a batch size of 50 in our evaluation, which is definitely small enough to fit in memory and also faster than using a batch size of 1. The entire evaluation process takes about a minute.

```
input_fn = lambda:create_tensorflow_dataset(
    'eval.tfrecords', 50, training=False)
regression_model.evaluate(input_fn)
```

Using the Estimator object (`regression_model`) to evaluate the model trained with 1.9M steps. The loss (MAE) on the evaluation set is about 3267.

Code for evaluating the regression model is shown below

```
class SalesModel(object):
    def __init__(self, hidden_layers):
        self.hidden_layers = hidden_layers

    def run_regression_eval(self, ckpt_dir):
        regression_model = self.create_regression_model(ckpt_dir)
        input_fn = lambda:create_tensorflow_dataset('eval.tfrecords', 50, training=False)
        return regression_model.evaluate(input_fn)

    def run_regression_training(self, ckpt_dir, batch_size, num_training_steps=None):
        regression_model = self.create_regression_model(ckpt_dir)
        input_fn = lambda:create_tensorflow_dataset('train.tfrecords', batch_size)
        regression_model.train(input_fn, steps=num_training_steps)

    def create_regression_model(self, ckpt_dir):
        config = tf.estimator.RunConfig(log_step_count_steps=5000)
        regression_model = tf.estimator.Estimator(
            self.regression_fn,
            config=config,
            model_dir=ckpt_dir)
        return regression_model

    def regression_fn(self, features, labels, mode, params):
        feature_columns = create_feature_columns()
        inputs = tf.feature_column.input_layer(features, feature_columns)
        batch_predictions = self.model_layers(inputs)
        predictions = tf.squeeze(batch_predictions)
```

```

if labels is not None:
    loss = tf.losses.absolute_difference(labels, predictions)

if mode == tf.estimator.ModeKeys.TRAIN:
    global_step = tf.train.get_or_create_global_step()
    adam = tf.train.AdamOptimizer()
    train_op = adam.minimize(
        loss, global_step=global_step)
    return tf.estimator.EstimatorSpec(mode, loss=loss, train_op=train_op)
if mode == tf.estimator.ModeKeys.EVAL:
    return tf.estimator.EstimatorSpec(mode, loss=loss)
if mode == tf.estimator.ModeKeys.PREDICT:
    prediction_info = {
        'predictions': batch_predictions
    }
    return tf.estimator.EstimatorSpec(mode, predictions=prediction_info)

def model_layers(self, inputs):
    layer = inputs
    for num_nodes in self.hidden_layers:
        layer = tf.layers.dense(layer, num_nodes,
            activation=tf.nn.relu)
    batch_predictions = tf.layers.dense(layer, 1)
    return batch_predictions

```