

Types of Locks: `std::lock_guard`

This lesson gives an introduction to locks and explains how `std::lock_guard` is used in C++.

WE'LL COVER THE FOLLOWING ^

- `std::lock_guard`

Locks take care of their resources following the RAII idiom. A lock automatically binds its mutex in the constructor and releases it in the destructor; this considerably reduces the risk of a deadlock because the runtime takes care of the mutex.

Locks are available in three different flavors: `std::lock_guard` for the simple use-cases; `std::unique_lock` for the advanced use-cases; `std::shared_lock` is available (with C++14) and can be used to implement reader-writer locks.

First, the simple use-case:

```
std::mutex m;  
m.lock();  
sharedVariable = getVar();  
m.unlock();
```



The mutex `m` ensures that access to the critical section `sharedVariable=getVar()` is sequential. Sequential, in this special case, means that each thread gains access to the critical section one after the other. This maintains a kind of total order in the system. The code is simple but prone to deadlocks. A deadlock appears if the critical section throws an exception or if the programmer simply forgets to unlock the mutex.

`std::lock_guard`

With `std::lock_guard` we can do this in a more elegant way:

With `std::lock_guard` we can do this in a more elegant way.

```
{  
    std::mutex m;  
    std::lock_guard<std::mutex> lockGuard(m);  
    sharedVariable= getVar();  
}
```

That was easy, but what's the story with the opening and closing brackets? The lifetime of `std::lock_guard` is limited by the curly **bracket**, which means that its lifetime ends when it passes the closing curly brackets. At exactly that point, the `std::lock_guard` destructor is called and - as we may have guessed - the mutex is released automatically. It is also released if `getVar()` in `sharedVariable = getVar()` throws an exception. The function scope and loop scope also limit the lifetime of an object.

i `std::scoped_lock` with C++17

With C++17, we get an `std::scoped_lock`. It's very similar to `std::unique_lock`, but `std::scoped_lock` can lock an arbitrary number of mutexes atomically. That being said, we have to keep two facts in mind:

1. If one of the current threads already owns the corresponding mutex and the mutex is not recursive, the behaviour is undefined.
2. We can only take ownership of the mutexes without locking them. In this case, we have to provide the `std::adopt_lock_t` flag to the constructor: `std::scoped_lock(std::adopt_lock_t, MutexTypes& ... m)`.

We can quite elegantly solve the previous deadlock by using an `std::scoped_lock`. We will discuss the resolution to the deadlock in the following section.

In the next lesson, we will see how `std::unique_lock` is stronger but more expensive than its counterpart `std::lock_guard`.

