

Types of Locks: `std::unique_lock`

This lesson gives an overview of `std::unique_lock` which is a type of lock used in C++.

WE'LL COVER THE FOLLOWING



- Features
- Methods
 - More on `lk.try_lock` and `lk.release` methods
 - How to solve deadlock with `std::unique_lock`?

Features

In addition to what's offered by an `std::lock_guard`, an `std::unique_lock` enables us to:

- create it without an associated mutex.
- create it without locking the associated mutex.
- explicitly and repeatedly set or release the lock of the associated mutex.
- move the mutex.
- try to lock the mutex.
- delay the lock on the associated mutex.

Methods

The following table shows the methods of an `std::unique_lock lk`.

Method	Description
<code>lk.lock()</code>	Locks the associated mutex.

<code>std::lock(lk1, lk2, ...)</code>	Atomically locks the arbitrary number of associated mutexes.
<code>lk.try_lock()</code> , and <code>lk.try_lock_for(relTime)</code> , and <code>lk.try_lock_until(absTime)</code>	Tries to lock the associated mutex.
<code>lk.release()</code>	Releases the mutex. The mutex remains locked.
<code>lk.swap(lk2)</code> and <code>std::swap(lk, lk2)</code>	Swaps the locks.
<code>lk.mutex()</code>	Returns a pointer to the associated mutex.
<code>lk.owns_lock()</code>	Checks if the lock has a mutex.

More on `lk.try_lock` and `lk.release` methods

`lk.try_lock_for(relTime)` needs a relative time duration;

`lk.try_lock_until(absTime)` needs an absolute time point.

`lk.try_lock` tries to lock the mutex and returns immediately. On success, it returns true, but otherwise, it's false. In contrast, the methods `lk.try_lock_for` and `lk.try_lock_until` block the release until the specified timeout occurs or the lock is acquired, whichever comes first. we should use a steady clock for our time constraint. A steady clock cannot be adjusted.

The method `lk.release()` returns the mutex; therefore, we have to unlock it manually.

How to solve deadlock with `std::unique_lock`?

Thanks to `std::unique_lock` it is quite easy to lock many mutexes in one

Thanks to `std::unique_lock`, it is quite easy to lock many mutexes in one atomic step; therefore, we can overcome deadlocks by locking mutexes in a different order. Remember the deadlock from the subsection [Issues of Mutexes?](#)

```
// deadlock.cpp

#include <iostream>
#include <chrono>
#include <mutex>
#include <thread>

struct CriticalData{
    std::mutex mut;
};

void deadLock(CriticalData& a, CriticalData& b){

    a.mut.lock();
    std::cout << "get the first mutex" << std::endl;
    std::this_thread::sleep_for(std::chrono::milliseconds(1));
    b.mut.lock();
    std::cout << "get the second mutex" << std::endl;
    // do something with a and b
    a.mut.unlock();
    b.mut.unlock();

}

int main(){

    CriticalData c1;
    CriticalData c2;

    std::thread t1([&]{deadLock(c1,c2);});
    std::thread t2([&]{deadLock(c2,c1);});

    t1.join();
    t2.join();

}
```



Let's solve this issue. The function `deadLock` has to lock its mutexes atomically and that's exactly what happens in the following example.

```
// deadlockResolved.cpp

#include <iostream>
#include <chrono>
#include <mutex>
#include <thread>
```

```

using namespace std;

struct CriticalData{
    mutex mut;
};

void deadLock(CriticalData& a, CriticalData& b){

    unique_lock<mutex> guard1(a.mut,defer_lock);
    cout << "Thread: " << this_thread::get_id() << " first mutex" << endl;

    this_thread::sleep_for(chrono::milliseconds(1));

    unique_lock<mutex> guard2(b.mut,defer_lock);
    cout << " Thread: " << this_thread::get_id() << " second mutex" << endl;

    cout << " Thread: " << this_thread::get_id() << " get both mutex" << endl;
    lock(guard1,guard2);
    // do something with a and b
}

int main(){

    cout << endl;

    CriticalData c1;
    CriticalData c2;

    thread t1([&]{deadLock(c1,c2);});
    thread t2([&]{deadLock(c2,c1);});

    t1.join();
    t2.join();

    cout << endl;

}

```



If we call the constructor of `std::unique_lock` with `std::defer_lock`, the underlying mutex will not be locked automatically. At this point (lines 16 and 21), the `std::unique_lock` is just the owner of the mutex. Thanks to the variadic template `std::lock`, the lock operation is performed in an atomic step (line 25). A variadic template is a template which can accept an arbitrary number of arguments. `std::lock` tries to get all locks in one atomic step, so it either gets all of them or none of them and retries until it succeeds.

In this example, `std::unique_lock` manages the lifetime of the resources and `std::lock` locks the associated mutex; we can also do it the other way around. In the first step the mutexes are locked, in the second `std::unique_lock`

manages the lifetime of resources. Here is an example of the second approach.

```
std::lock(a.mut, b.mut);  
std::lock_guard<std::mutex> guard1(a.mut, std::adopt_lock);  
std::lock_guard<std::mutex> guard2(b.mut, std::adopt_lock);
```

Let us see this approach in action:

```
// deadlockResolved.cpp  
  
#include <iostream>  
#include <chrono>  
#include <mutex>  
#include <thread>  
  
using namespace std;  
  
struct CriticalData{  
    mutex mut;  
};  
  
void deadLock(CriticalData& a, CriticalData& b){  
  
    lock_guard<std::mutex> guard1(a.mut, std::adopt_lock);  
    cout << "Thread: " << this_thread::get_id() << " first mutex" << endl;  
  
    this_thread::sleep_for(chrono::milliseconds(1));  
  
    lock_guard<std::mutex> guard2(b.mut, std::adopt_lock);  
    cout << " Thread: " << this_thread::get_id() << " second mutex" << endl;  
  
    cout << "      Thread: " << this_thread::get_id() << " get both mutex" << endl;  
    lock(a.mut, b.mut);  
    // do something with a and b  
}  
  
int main(){  
  
    cout << endl;  
  
    CriticalData c1;  
    CriticalData c2;  
  
    thread t1([&]{deadLock(c1,c2);});  
    thread t2([&]{deadLock(c2,c1);});  
  
    t1.join();  
    t2.join();  
  
    cout << endl;  
  
}
```



i Resolving the deadlock with an `std::scoped_lock`

With C++17, the resolution to the deadlock becomes quite easy. We get the `std::scoped_lock` that can lock an arbitrary number of mutexes atomically - so as long as we only have to use an `std::lock_guard` instead of the `std::lock` call. That's all. Here is the modified function `deadlock`.

```
// deadlockResolvedScopedLock.cpp
void deadLock(CriticalData& a, CriticalData& b){

    cout << "Thread: " << this_thread::get_id() << " first mutex" << endl;
    this_thread::sleep_for(chrono::milliseconds(1));
    cout << "  Thread: " << this_thread::get_id() << " second mutex" << endl;
    cout << "    Thread: " << this_thread::get_id() << " get both mutex" << endl;

    std::scoped_lock(a.mut, b.mut);
    // do something with a and b
}
```



With C++14, C++ adds support for `std::shared_lock`; let's see this in the next lesson.