

Memory layout

This lesson explains the memory layout using NumPy.

WE'LL COVER THE FOLLOWING



- Layouts
 - Item layout
 - Example 1
 - Example 2
 - Flattened item layout
 - Example 1
 - Example 2
 - Memory layout (C order, big endian)
 - Example 1
 - Example 2

The [NumPy documentation](#) defines the ndarray class very clearly:

*An instance of class **ndarray** consists of a contiguous one-dimensional segment of computer memory (owned by the array, or by some other object), combined with an indexing scheme that maps N integers into the location of an item in the block.*

Said differently, an array is mostly a contiguous block of memory whose parts can be accessed using an indexing scheme. Such indexing scheme is in turn defined by a [shape](#) and a [data type](#) and this is precisely what is needed when you define a new array:

```
import numpy as np
Z = np.arange(9).reshape(3,3).astype(np.int16)
print(Z)
```



```
print(Z.itemsize)#returns size of Z in bytes
print(Z.shape)# returns the x dimension and y dimension of Z
print(Z.ndim)# dimension in Z i.e (2 in this case) since the array is 2D
```



Here, we know that itemsize is 2 bytes (`int16`), the shape is (3,3) and the number of dimensions is 2.

To calculate the dimension we can also use `len(Z.shape)` .

Furthermore, we can deduce the `strides` of the array that define the number of bytes to step in each dimension when traversing the array.

```
import numpy as np
Z = np.arange(9).reshape(3,3).astype(np.int16)
stride = Z.shape[1]*Z.itemsize, Z.itemsize # store stride of Z
print("Stride(as np.int16):",stride)
print("Z.strides:",Z.strides)
Z = np.arange(9).reshape(3,3).astype(np.int32)
stride= Z.shape[1]*Z.itemsize, Z.itemsize #stores stride of Z
print("Stride(as np.int32):",stride)
print("Z.strides:",Z.strides)
```



Here in this example, we have to skip 2 bytes (1 value) to move to the next column, but 6 bytes (3 values) to get to the same position in the next row. As such, the **strides** for the array `Z` will be **(6, 2)**.

With all this information, we know how to access a specific item (designed by an index tuple) and more precisely, how to compute the start and end offsets:

```
import numpy as np
Z = np.arange(9).reshape(3,3).astype(np.int16)
offset_start = 0
for i in range(Z.ndim):
    offset_start += Z.strides[i] * i #compute the start offset of Z
    offset_end = offset_start + Z.itemsize #compute the end offset of Z

print("Starting offset:", offset_start)
print("Ending offset:", offset_end)
```



Let's see if this is correct using the `tobytes` conversion method that construct Python bytes containing the raw data bytes in the array:

```
import numpy as np
Z = np.arange(9).reshape(3, 3).astype(np.int16)
index = 1, 1
print(Z[index].tobytes())
#b'\x04\x00'
offset_start = 0
for i in range(Z.ndim):
    offset_start += Z.strides[i] * index[i]
    offset_end = offset_start + Z.itemsize

print(Z.tobytes()[offset_start:offset_end])#b'\x04\x00'
```

Layouts

This array can be actually considered from different perspectives (i.e. layouts):

Item layout

Consider item layout as a 2-Dimensional Matrix with x rows and y columns.

Example 1


Consider the following example:

```
Z = np.arange(9).reshape(3, 3)
```

Here takes 9 values from 0-8 and is reshaped in 2-D matrix format having dimensions (3 * 3).

Z has the following item layout:

shape[1]
(=3)



```

shape[0] | 3 | 4 | 5 | len(Z)
(=3)     |---|---|---| (=3)
         | 6 | 7 | 8 |
         |---|---|---|

```

Example 2

Consider the following example

```

Z = np.arange(9).reshape(3,3)
V= Z[:,2,:2])

```

Here Z takes up 9 values from 0-8 and arranges them in a 3 * 3 matrix V takes values from corners from the grid. i.e V has 4 values.

V has the following item layout:

```

              shape[1]
              (=2)
              |---|
shape[0]      | 0 |   | 2 |
(=2)          |---|---|
              |   |   |   |
              | 6 |   | 8 |
              |---|---|
len(V)        | 0 | 2 |
(=2)          |---|
              | 6 | 8 |
              |---|

```

Flattened item layout

Consider flattened item layout as 1-Dimensional Matrix with 1 row and n columns.

Example 1

Consider the following example:

```

Z = np.arange(9)

```

It makes 9 indices in the computer's memory and places values from 0 to 8.

Z has the following flattened item layout:

Z has the following flattened item layout.



`Z.size`
(=9)

Example 2

Consider the following example

```
Z = np.arange(9).reshape(3,3).astype(np.int16)
V = Z[::2,::2]
V=V.reshape(1,4)
```

Here Z takes up 9 values from 0-8 and arranges them in a 3 * 3 matrix V takes values from corners from the grid. i.e V has 4 values. V has the following flattened item layout:



`V.size`
(=4)

Memory layout (C order, big endian)

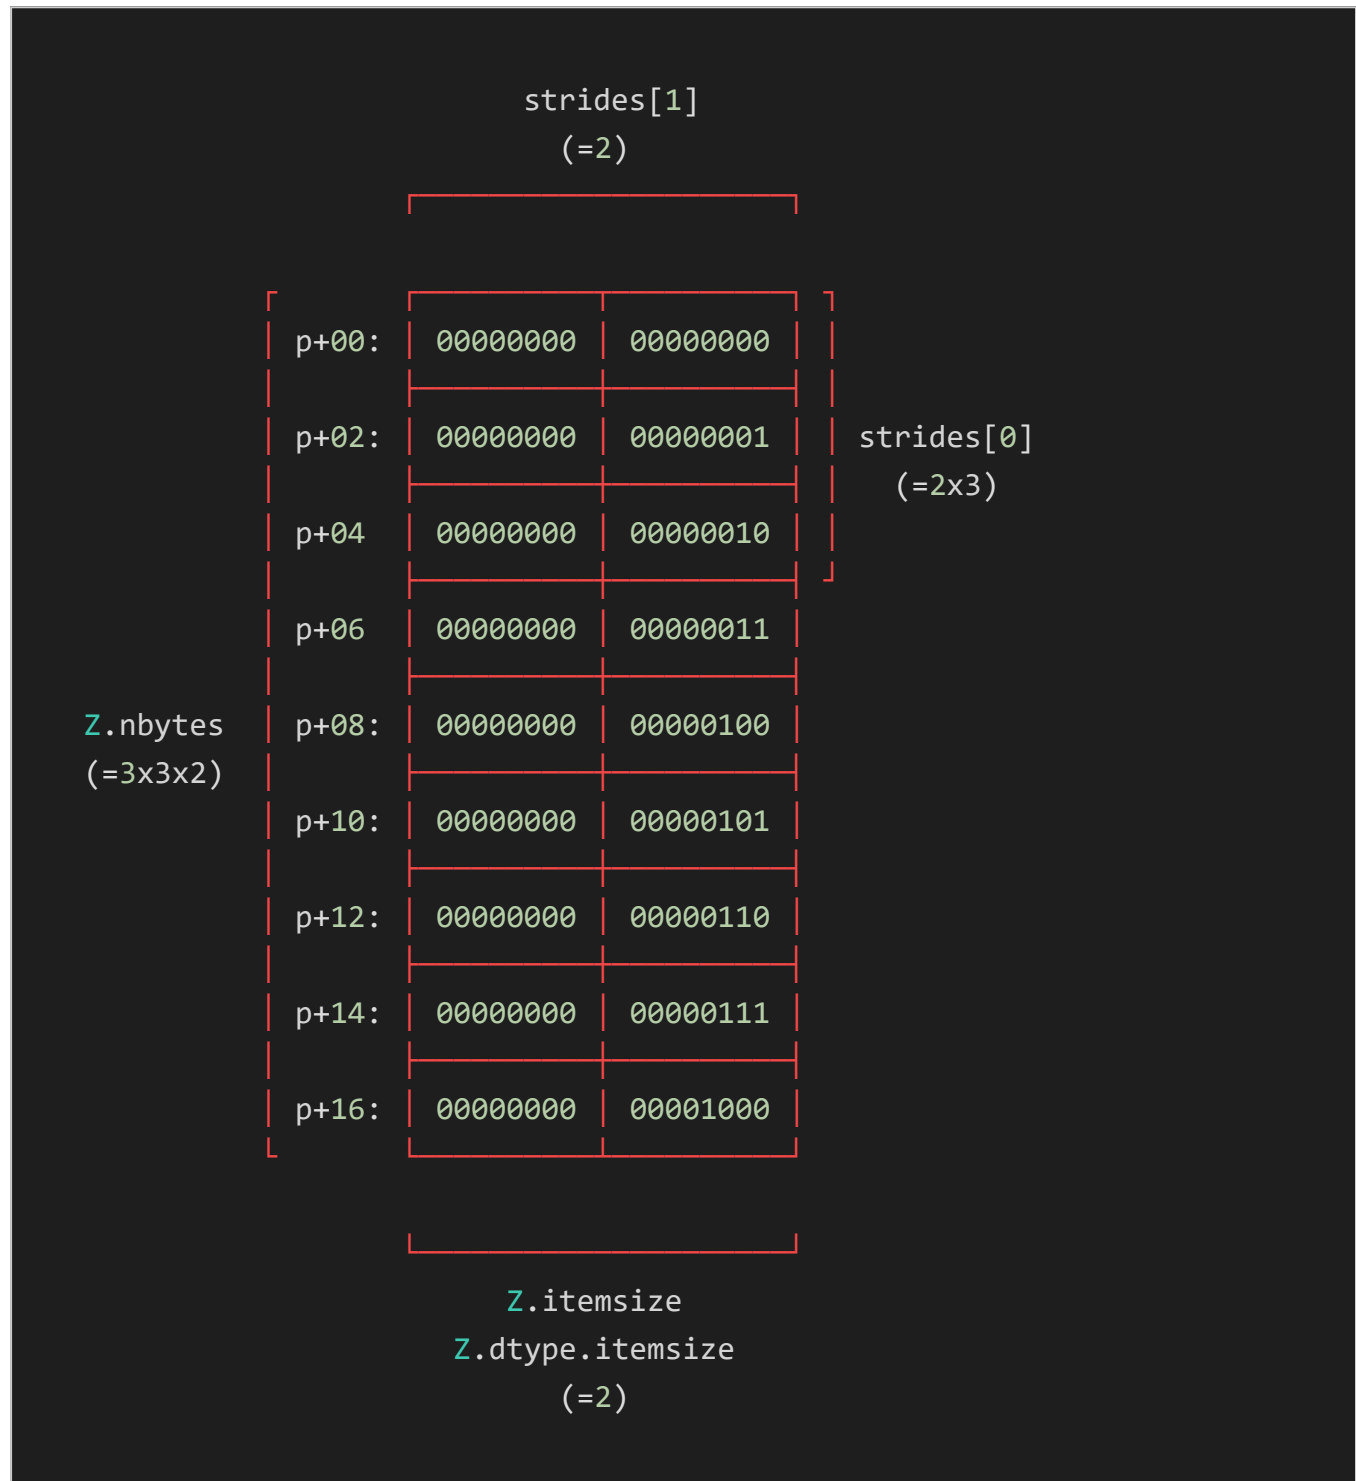
Consider memory layout with rows equal to the number of bytes and columns equal to the number of bytes divided by 8 (i.e Z.itemsize).

Example 1

Consider the following example:

```
Z = np.arange(9).reshape(3, 3).astype(np.int16)
```

Here, the number of rows is 16 and the number of columns is 2. The total number of bytes is $3 * 3 * 2$ where $3 * 3$ is the size of the grid and each cell takes 2 bytes, so total 18 bytes. Z has the following memory layout:



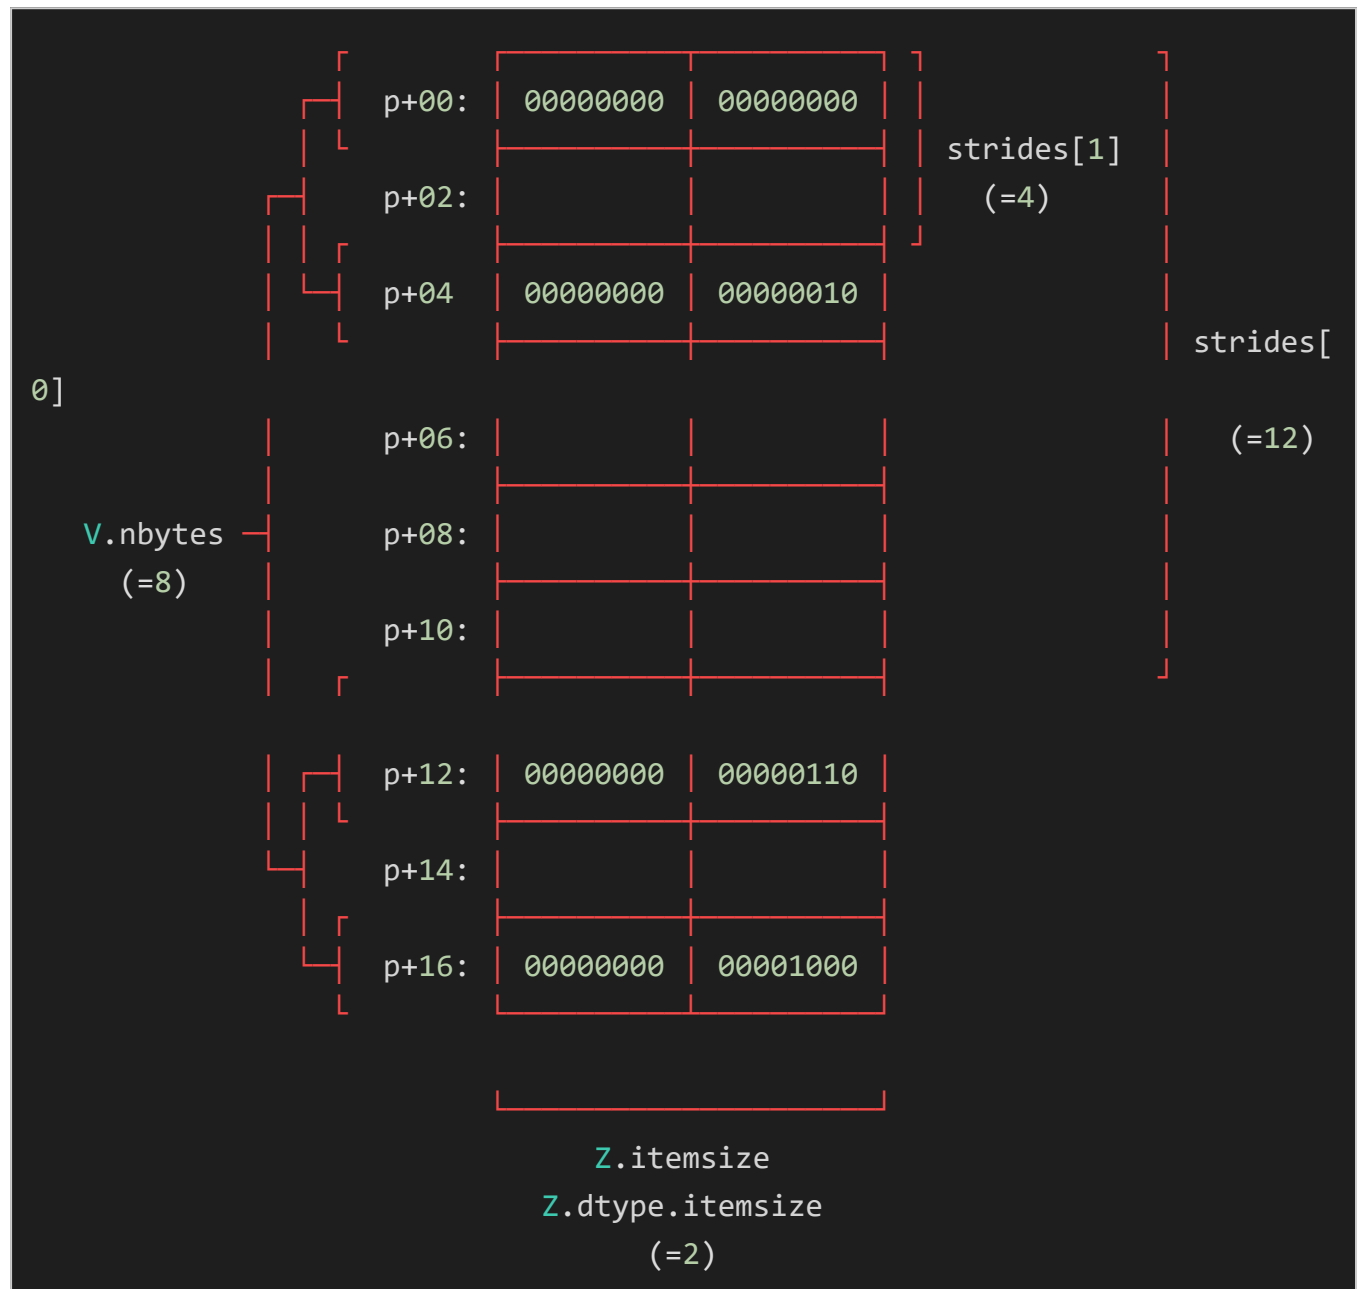
Example 2

Consider the following example

```
Z = np.arange(9).reshape(3,3).astype(np.int16)
V = Z[:, :2, ::2]
```

Here we take a slice of **Z**, the result is a view of the base array **Z**. In the

memory layout below, since the array takes up only 4 value. and each value is 2 bytes so the total bytes are $2 * 4 = 8$ bytes. V has the following memory layout:



Such a view is specified using a shape, a dtype and strides because strides cannot be deduced anymore from the dtype and shape only.

Solve this Quiz!

1

What is the output of the following code?

```
Z = np.arange(9).reshape(3,3).astype(np.int32)
print(Z.itemsize())
```

```
printf("%d\n", *p);
```

COMPLETED 0%

1 of 2



Now, that we have viewed memory layouts, we'll look at views and copies in the next lesson.