

Bayesian Regression

Learn about Bayesian regression techniques.

Chapter Goals:

- Learn about Bayesian regression techniques

A. Bayesian techniques

So far, we've discussed hyperparameter optimization through cross-validation. Another way to optimize the hyperparameters of a regularized regression model is with [Bayesian](#) techniques.

In Bayesian statistics, the main idea is to make certain assumptions about the probability distributions of a model's parameters *before* being fitted on data. These initial distribution assumptions are called *priors* for the model's parameters.

In a Bayesian ridge regression model, there are two hyperparameters to optimize: α and λ . The α hyperparameter serves the same exact purpose as it does for regular ridge regression; namely, it acts as a scaling factor for the penalty term.

The λ hyperparameter acts as the [precision](#) of the model's weights. Basically, the smaller the λ value, the greater the variance between the individual weight values.

B. Hyperparameter priors

Both the α and λ hyperparameters have [gamma distribution](#) priors, meaning we assume both values come from a gamma probability distribution.

There's no need to know the specifics of a gamma distribution, other than the fact that it's a probability distribution defined by a [shape parameter](#) and [scale parameter](#).

Specifically, the α hyperparameter has prior:

specifically, the α hyperparameter has prior:

$$\Gamma(\alpha_1, \alpha_2)$$

and the λ hyperparameter has prior:

$$\Gamma(\lambda_1, \lambda_2)$$

where $\Gamma(k, \theta)$ represents a gamma distribution with shape parameter k and scale parameter θ .

C. Tuning the model

When finding the optimal weight settings of a Bayesian ridge regression model for an input dataset, we also concurrently optimize the α and λ hyperparameters based on their prior distributions and the input data.

This can all be done with the `BayesianRidge` object (part of the `linear_model` module). Like all the previous regression objects, this one can be initialized with no required arguments.

```
# predefined dataset from previous chapter
print('Data shape: {}'.format(data.shape))
print('Labels shape: {}'.format(labels.shape))

from sklearn import linear_model
reg = linear_model.BayesianRidge()
reg.fit(data, labels)
print('Coefficients: {}'.format(repr(reg.coef_)))
print('Intercept: {}'.format(reg.intercept_))
print('R2: {}'.format(reg.score(data, labels)))
print('Alpha: {}'.format(reg.alpha_))
print('Lambda: {}'.format(reg.lambda_))
```



We can manually specify the α_1 and α_2 gamma parameters for α with the `alpha_1` and `alpha_2` keyword arguments when initializing `BayesianRidge`. Similarly, we can manually set λ_1 and λ_2 with the `lambda_1` and `lambda_2` keyword arguments. The default value for each of the four gamma parameters is 10^{-6} .

Time to Code!

The coding exercise in this chapter uses the `BayesianRidge` object of the `linear_model` module (imported in backend) to complete the `bayes_ridge` function.

The function will fit a Bayesian ridge regression model to the input data and labels.

Set `reg` equal to `linear_model.BayesianRidge`, initialized with no input arguments.

Call `reg.fit` with `data` and `labels` as the two input arguments. Then return `reg`.

```
def bayes_ridge(data, labels):  
    # CODE HERE  
    pass
```

