

Calculate Sum of a Vector: Conclusion

This lesson concludes all the methods used to solve the problem of calculating the sum of a vector in C++.

WE'LL COVER THE FOLLOWING



- Single Threaded
- Multithreading with a Shared Variable
- Thread-local Summation

Let's conclude what we learned from this chapter:

Single Threaded


The range-based for loop and the STL algorithm `std::accumulate` are in the same performance range. This observation holds for the most optimized version. In the optimized version, the compiler uses the optimized version vectorized `SIMD` instruction (SSE or AVX) for the summation case; therefore, the loop counter will be increased by 2 (SSE) or 4 (AVX).

Multithreading with a Shared Variable

The usage of a shared variable for the summation variable makes one point clear: synchronization is very expensive and should be avoided as much as possible. Although I used an atomic variable and even broke the sequential consistency, the four threads are 100 times slower than one thread. From a performance perspective, minimizing expensive synchronization has to be our first goal.

Thread-local Summation

The thread-local summation is only two times faster than the single-threaded range-based for loop or `std::accumulate`; that holds, even though each of the four threads can work independently. That surprised me because I was expecting a nearly fourfold improvement, and it surprised me even more that my four cores were not fully utilized.



The reason is simple: the cores can't get the data fast enough from the memory. To that point, the execution is **memory bound** - i.e. it slows down the cores.