

# CppMem: Atomics with an Acquire-Release Semantic

This lesson gives an overview of atomics with acquire-release semantic used in the context of CppMem.

## WE'LL COVER THE FOLLOWING



- Explanation
- CppMem
- Possible Executions
  - Execution for (y = 0, x = 0)
  - Execution for (y = 0, x = 2000)
  - Execution for (y = 11, x = 2000)

The synchronization in the [acquire-release semantic](#) takes place between atomic operations on the same atomic. This is in contrast to the sequential consistency where we have synchronization between threads. Due to this fact, the acquire-release semantic is more lightweight and, therefore, faster.

Here is the program with acquire-release semantic:

```
// ongoingOptimisationAcquireRelease.cpp

#include <atomic>
#include <iostream>
#include <thread>

std::atomic<int> x{0};
std::atomic<int> y{0};

void writing(){
    x.store(2000, std::memory_order_relaxed);
    y.store(11, std::memory_order_release);
}

void reading(){
    std::cout << y.load(std::memory_order_acquire) << " ";
    std::cout << x.load(std::memory_order_relaxed) << std::endl;
}

int main(){
```



```
std::thread thread1(writing);
std::thread thread2(reading);
thread1.join();

thread2.join();
};
```



On first glance you will notice that all operations are atomic, so the program is well-defined. But the second glance shows more; the atomic operations on `y` are attached with the flag `std::memory_order_release` (line 12) and `std::memory_order_acquire` (line 16). In contrast to that, the atomic operations on `x` are annotated with `std::memory_order_relaxed` (lines 11 and 17), so there are no synchronizations and ordering constraints for `x`. The answer to the possible values for `x` and `y` can only be given by `y`.

It holds:

- `y.store(11, std::memory_order_release)` **synchronizes-with** `y.load(std::memory_order_acquire)`
- `x.store(2000, std::memory_order_relaxed)` **is visible before** `y.store(11, std::memory_order_release)`
- `y.load(std::memory_order_acquire)` **is visible before** `x.load(std::memory_order_relaxed)`

## Explanation #

I will elaborate a little bit more on these three statements. The key idea is that the store of `y` in line 12 synchronizes with a load of `y` in line 16. This is due to the fact that the operations take place on the same atomic and they use the acquire-release semantic. `y` uses `std::memory_order_release` in line 12 and `std::memory_order_acquire` in line 16. The pairwise operation on `y` has another very interesting property. They establish a kind of barrier relative to `y`; so, `x.store(2000, std::memory_order_relaxed)` cannot be executed after `y.store(std::memory_order_release)` and `x.load()` cannot be executed before `y.load()`.

The reasoning in the case of the acquire-release semantic is more sophisticated than in the case of the previous sequential consistency, but the

possible values for `x` and `y` are the same; Only the combination `y == 11` and `x == 0` is not possible.

There are three different interleavings of the threads possible, which produce the three different combinations of the values `x` and `y`.

- `thread1` will be executed before `thread2`.
- `thread2` will be executed before `thread1`.
- `thread1` executes `x.store(2000)` before `thread2` will be executed.

To make a long story short, here are all possible values for `x` and `y`.

| y  | x    | Values possible? |
|----|------|------------------|
| 0  | 0    | Yes              |
| 11 | 0    |                  |
| 0  | 2000 | Yes              |
| 11 | 2000 | Yes              |

Once more, let's verify our thinking with CppMem.

## CppMem #

Here is the corresponding program:

```
int main(){
    atomic_int x = 0;
    atomic_int y = 0;
    {{{ {
        x.store(2000, memory_order_relaxed);
        y.store(11, memory_order_release);
    }
    }
```



```

    }
    }
    {
        y.load(memory_order_acquire);
        x.load(memory_order_relaxed);
    }
}
}


```

We already know that all results are possible except for  $(y = 11, x = 0)$ .

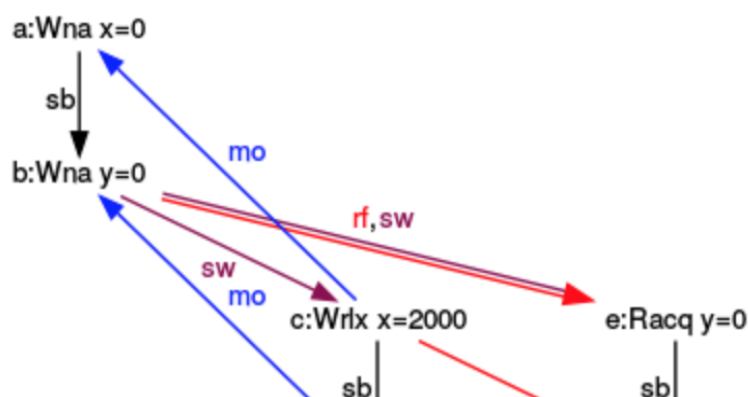
## Possible Executions #


I will only refer to the three graphs with consistent execution. The graphs show that there is an acquire-release semantic between the store-release of  $y$  and the load-acquire operation of  $y$ . It will not make any difference if the reading of  $y$  (**rf**) takes places in the main thread or in a separate thread. The graphs show the synchronizes-with relation using a **sw** annotated arrow.

### Execution for $(y = 0, x = 0)$ #




### Execution for $(y = 0, x = 2000)$ #





Execution for ( $y = 11, x = 2000$ ) #



x does not have to be atomic. This was my first and wrong assumption; see why in the next lesson. :)