

# The Two Flavors of Transactional Memory

This lesson explains synchronized and atomic blocks in transactional memory.

## WE'LL COVER THE FOLLOWING



- Synchronized & Atomic Blocks
- Synchronized Blocks
- Atomic Blocks
- `transaction_safe` versus `transaction_unsafe` Code

C++ supports transactional memory in two flavors: synchronized blocks and atomic blocks.

## Synchronized & Atomic Blocks #

Up to now, I only wrote about transactions. Now, I will write about synchronized blocks and atomic blocks. Both can be encapsulated in each other; specifically, synchronized blocks are not atomic blocks because they can execute transaction-unsafe. An example would be code like the output to the console which can not be undone. For this reason, synchronized blocks are often called relaxed blocks.

## Synchronized Blocks #

Synchronized blocks behave like they are protected by a global lock, i.e. This means that all synchronized blocks follow a total order. In particular: all changes to a synchronized block are available in the next synchronized block. There is a *synchronizes-with* relation between the synchronized blocks

because of the commit of the transaction *synchronizes-with* the next start of a transaction. Synchronized blocks cannot cause a deadlock because they create a total order. While a classical lock protects a memory area, a global lock of a synchronized block protects the total program. This is the reason the following program is *well-defined*:

```
// synchronized.cpp

#include <iostream>
#include <vector>
#include <thread>

int i= 0;

void increment(){
    synchronized{
        std::cout << ++i << " ,";
    }
}

int main(){

    std::cout << std::endl;

    std::vector<std::thread> vecSyn(10);
    for(auto& thr: vecSyn)
        thr = std::thread([]{ for(int n = 0; n < 10; ++n) increment(); });
    for(auto& thr: vecSyn) thr.join();

    std::cout << "\n\n";

}
```

Although the variable `i` in line 7 is a global variable and the operations in the synchronized block are transaction-unsafe, the program is well-defined. 10 threads concurrently invoke the function `increment` (line 21) ten times, incrementing the variable `i` in line 11. The access to `i` and `std::cout` happens in total order. This is the characteristic of the synchronized block. Afterwards, the program returns the expected result; the values for `i` are written in an increasing sequence, separated by a comma.

What about data races? You can have them with synchronized blocks. A small modification of the source code is sufficient to introduce a [data race](#)

```
// nonsynchronized.cpp
```

```
#include <chrono>
#include <iostream>
#include <vector>
```

```

#include <thread>

using namespace std::chrono_literals;

using namespace std;

int i= 0;

void increment(){
    synchronized{
        cout << ++i << " ,";
        this_thread::sleep_for(1ns);
    }
}

int main(){

    cout << endl;

    vector<thread> vecSyn(10);
    vector<thread> vecUnsyn(10);

    for(auto& thr: vecSyn)
        thr = thread([]{ for(int n = 0; n < 10; ++n) increment(); });
    for(auto& thr: vecUnsyn)
        thr = thread([]{ for(int n = 0; n < 10; ++n) cout << ++i << " ,"; });

    for(auto& thr: vecSyn) thr.join();
    for(auto& thr: vecUnsyn) thr.join();

    cout << "\n\n";

}

```

To observe the data race, I let the synchronized block sleep for a nanosecond (line 16). At the same time I access the output stream `std::cout` without a synchronized block (line 30). In total, 20 threads increment the global variable `i` - half of them without synchronization. The C++11 standard guarantees that the characters will be written atomically; that is not an issue. What is worse is that the variable `i` is written by at least 2 threads. This is a data race, hence, the program has undefined behavior. The total order of synchronized blocks also holds for atomic blocks.

## Atomic Blocks #

You can execute transaction-unsafe code in a synchronized block, but not in an atomic block. Atomic blocks are available in three forms: `atomic_noexcept`, `atomic_commit`, and `atomic_cancel`. The three suffixes `_noexcept`, `_commit`, and `_cancel` define how an atomic block should manage an exception.

`atomic_noexcept`: If an exception is thrown, `std::abort` will be called and the program aborts.

`atomic_cancel`: In the default case, `std::abort` is called. This will not hold if a transaction-safe exception is thrown that is responsible for ending the transaction. In this case the transaction will be canceled, put to its initial state, and the exception will be thrown.

`atomic_commit`: If an exception is thrown, the transaction will be committed.

Transaction-safe exceptions are: `std::bad_alloc`, `std::bad_array_length`, `std::bad_array_new_length`, `std::bad_cast`, `std::bad_typeid`, `std::bad_exception`, `std::exception`, and all exceptions are derived from one of these.

## `transaction_safe` versus `transaction_unsafe` Code #

You can declare a function as `transaction_safe` or attach the `transaction_unsafe` attribute to it.

```
int transactionSafeFunction() transaction_safe;

[[transaction_unsafe]] int transactionUnsafeFunction();
```



`transaction_safe` belongs to the type of the function. What does `transaction_safe` mean? A `transaction_safe` function is, according to the proposal [N4265](#), a function that has a `transaction_safe` definition. This holds true if the following properties *do not* apply to its definition:

- It has a `volatile` parameter or a `volatile` variable.
- It has `transaction_unsafe` statements.
- If the function uses a constructor or destructor of a class in its body that has a `volatile` non-static member.

Of course this definition of `transaction_safe` is not sufficient because it uses the term `transaction_unsafe`. You can read the proposal [N4265](#) for the details.

