

# Data Races

This lesson gives an overview of data race problems which might occur during the implementation of concurrency in C++.

A data race is a situation in which at least two threads access a shared variable at the same time. Within that, at least one thread tries to modify the variable.

If your program has a data race, it will have undefined behavior. This means all outcomes are possible and, therefore, reasoning about the program makes no sense anymore. Let me show you a program with a data race.

```
// addMoney.cpp

#include <functional>
#include <iostream>
#include <thread>
#include <vector>

struct Account{
    int balance{100};
};

void addMoney(Account& to, int amount){
    to.balance += amount;
}

int main(){

    std::cout << std::endl;

    Account account;

    std::vector<std::thread> vecThreads(100);

    for (auto& thr: vecThreads) thr = std::thread(addMoney, std::ref(account), 50);

    for (auto& thr: vecThreads) thr.join();

    std::cout << "account.balance: " << account.balance << std::endl;

    std::cout << std::endl;

}
```



100 threads are adding 50 euros (line 25) to the same account (line 20). They use the function `addMoney`. The key observation is that the writing to the account is done without synchronization, therefore we have a data race and the result is not valid. This is undefined behavior and the final balance (line 30) differs between 5000 and 5100 euro.