

Multithreading: Condition Variables

This lesson elucidates best practices involving condition variables such as `notify_one` and `notify_all` in multithreaded applications in C++.

WE'LL COVER THE FOLLOWING



- Condition Variables
 - Don't use condition variables without a predicate
 - Use Promises and Futures instead of Condition Variables
- Promises and Futures
 - If possible, go for `std::async`

Condition Variables

Synchronizing threads via notifications is a simple concept, but **condition variables** make this task really challenging - mostly because the condition variables have no state.

- If a condition variable gets a notification, it may be the wrong one **spurious wakeup**.
- If a condition variable gets its notification before it was ready, the notification will be lost **lost wakeup**.

Don't use condition variables without a predicate

Using a condition variable without a predicate is a **race condition**.

```
// conditionVariableLostWakeup.cpp
```

```
#include <condition_variable>
#include <mutex>
#include <thread>
```

```
std::mutex mutex ;
```



```

std::condition_variable condVar;

void waitingForWork(){
    std::unique_lock<std::mutex> lck(mutex_);
    condVar.wait(lck);
    // do the work
}

void setDataReady(){
    condVar.notify_one();
}

int main(){

    std::thread t1(setDataReady);
    std::thread t2(waitingForWork);

    t1.join();
    t2.join();

}

```



If the thread **t1** runs before the thread **t2**, you will get a deadlock. **t1** will send its notification before **t2** can accept it, and the notification is lost. This will happen very often because thread **t1** starts before thread **t2**, and thread **t1** has less work to perform.

Adding a bool variable **dataReady** to the workflow will solve this issue. **dataReady** will also protect against a **spurious wakeup**, as the waiting thread first checks if the notification was from the right thread.

```

// conditioVarialbleLostWakeupSolved.cpp

#include <condition_variable>
#include <mutex>
#include <thread>

std::mutex mutex_;
std::condition_variable condVar;

bool dataReady{false};

void waitingForWork(){
    std::unique_lock<std::mutex> lck(mutex_);
    condVar.wait(lck, []{ return dataReady; });
    // do the work
}

void setDataReady(){
    {
        std::lock_guard<std::mutex> lck(mutex_);

```



```

        dataReady = true;
    }
    condVar.notify_one();
}

int main(){

    std::thread t1(waitingForWork);
    std::thread t2(setDataReady);

    t1.join();
    t2.join();

}

```



Use Promises and Futures instead of Condition Variables

For one-time notifications, [promises and futures](#) are the better choice. The workflow of the previous program `conditioVarialbleLostWakeupSolved.cpp` can directly be implemented with a promise and a future.

```

// notificationWithPromiseAndFuture.cpp

#include <future>
#include <utility>

void waitingForWork(std::future<void>&& fut){
    fut.wait();
    // do the work
}

void setDataReady(std::promise<void>&& prom){
    prom.set_value();
}

int main(){

    std::promise<void> sendReady;
    auto fut = sendReady.get_future();

    std::thread t1(waitingForWork, std::move(fut));
    std::thread t2(setDataReady, std::move(sendReady));

    t1.join();
    t2.join();

}

```



The workflow is reduced to its bare minimum. The promise `prom.set_value()` sends the notification the future `fut.wait()` is waiting for. The program needs no mutexes and locks because there is no critical section. Because no [lost wakeup](#) or [spurious wakeup](#) can happen, a predicate is also not necessary.

If your workflow requires that you use a condition variable many times, then a promise and a future pair is no alternative.

Promises and Futures

`std::async` can often be used as an easy-to-use replacement for threads or condition variables.

If possible, go for `std::async` #

If possible, you should go for `std::async` to execute an asynchronous task.

```
auto fut = std::async([]{ return 2000 + 11; });
// some time passes
std::cout << "fut.get(): " << fut.get() << std::endl;
```



Let's see this in action:

```
#include <future>
#include <thread>
#include <iostream>
#include <chrono>

int main(){

    std::cout << std::endl;

    auto fut= std::async([]{ return 2000 + 11; });
    std::this_thread::sleep_for (std::chrono::seconds(2)); //Work for 2 seconds
    std::cout << "fut.get(): " << fut.get() << std::endl;

    std::cout << std::endl;

}
```



By invoking `auto fut = std::async([]{ return 2000 + 11; })`, you say to the C++ runtime: “Run my job”. I don’t care if it will be executed immediately, if it

will run on the same thread, if it will run on a thread pool, or if it will run on a GPU; You are only interested in picking up the result in the future: `fut.get()`.

From a conceptional view, a thread is just an implementation detail for running your job. You only specify *what* should be done and not *how* it should be done.