

# Sleep and Wait

This lesson gives a brief introduction to sleep/wait and its usage in C++ with the help of interactive examples.

WE'LL COVER THE FOLLOWING ^

- Conventions
- Various waiting strategies

One important feature that multithreading components such as threads, locks, condition variables, and futures have in common is the notion of time.

## Conventions #

The methods for handling time in multithreading programs follow a simple convention: Methods ending with `_for` have to be parametrized by a [time duration](#); methods ending with `_until` by a [time point](#). Here is a concise overview of the methods that deal with sleeping, blocking, and waiting:

Multithreading Component	<code>_until</code>	<code>_for</code>
<code>std::thread th</code>	<code>th.sleep_until(in2min n)</code>	<code>th.sleep_for(2s)</code>
<code>std::unique_lock lk</code>	<code>lk.try_lock_until(in 2min)</code>	<code>lk.try_lock(2s)</code>
<code>std::condition_variable cv</code>	<code>cv.wait_until(in2min )</code>	<code>cv.wait_for(2s)</code>

<code>std::future fu</code>	<code>fu.wait_until(in2min )</code>	<code>fu.wait_for(2s)</code>
<code>std::shared_future shFu</code>	<code>shFu.wait(in2min)</code>	<code>shFu.wait_for(2s)</code>

`in2min` stands for a time 2 minutes in the future; `2s` is a time duration of 2 seconds. Although I use `auto` in the initialization of the time point `in2min`, the following is still verbose:

```
auto in2min= std::chrono::steady_clock::now() + std::chrono::minutes(2);
```



*Time literals* from C++14 come to our rescue when using typical time durations, e.g. `2s` stands for 2 seconds. Let's look at different waiting strategies.

## Various waiting strategies #

The main idea of the following program is that the promise provides its result for four shared futures. That is possible because more than one `shared_future` can wait for the notification of the same promise. Each future has a different waiting strategy. Both the promise and every future will be executed in different threads. For simplicity reasons, I will only speak about a waiting thread in the rest of this subsection, although it will be the corresponding future that is waiting. Below are the details of the [promises and the futures](#).

Here are the strategies for the four waiting threads:

- `consumeThread2`: waits up to 20 seconds for the result of the promise.
- `consumeThread3`: asks the promise for the result and goes back to sleep for 700 milliseconds.
- `consumeThread4`: asks the promise for the result and goes back to sleep. Its sleep duration starts with 1 millisecond and doubles each time.

Here is the program:



```
// sleepAndWait.cpp
```

```
#include <utility>
#include <iostream>
#include <future>
#include <thread>
#include <utility>

using namespace std;
using namespace std::chrono;

mutex coutMutex;

long double getDifference(const steady_clock::time_point& tp1,
                          const steady_clock::time_point& tp2){
    const auto diff= tp2 - tp1;
    const auto res= duration <long double, milli> (diff).count();
    return res;
}

void producer(promise<int>&& prom){
    cout << "PRODUCING THE VALUE 2011\n\n";
    this_thread::sleep_for(seconds(5));
    prom.set_value(2011);
}

void consumer(shared_future<int> fut,
              steady_clock::duration dur){
    const auto start = steady_clock::now();
    future_status status= fut.wait_until(steady_clock::now() + dur);
    if ( status == future_status::ready ){
        lock_guard<mutex> lockCout(coutMutex);
        cout << this_thread::get_id() << " ready => Result: " << fut.get()
              << endl;
    }
    else{
        lock_guard<mutex> lockCout(coutMutex);
        cout << this_thread::get_id() << " stopped waiting." << endl;
    }
    const auto end= steady_clock::now();
    lock_guard<mutex> lockCout(coutMutex);
    cout << this_thread::get_id() << " waiting time: "
          << getDifference(start,end) << " ms" << endl;
}

void consumePeriodically(shared_future<int> fut){
    const auto start = steady_clock::now();
    future_status status;
    do {
        this_thread::sleep_for(milliseconds(700));
        status = fut.wait_for(seconds(0));
        if (status == future_status::timeout) {
            lock_guard<mutex> lockCout(coutMutex);
            cout << "      " << this_thread::get_id()
                  << " still waiting." << endl;
        }
        if (status == future_status::ready) {
            lock_guard<mutex> lockCout(coutMutex);
            cout << "      " << this_thread::get_id()
                  << " waiting done => Result: " << fut.get() << endl;
        }
    }
```

```

    } while (status != future_status::ready);
    const auto end= steady_clock::now();
    lock_guard<mutex> lockCout(coutMutex);

    cout << "          " << this_thread::get_id() << " waiting time: "
          << getDifference(start,end) << " ms" << endl;
}

void consumeWithBackoff(shared_future<int> fut){
    const auto start = steady_clock::now();
    future_status status;
    auto dur= milliseconds(1);
    do {
        this_thread::sleep_for(dur);
        status = fut.wait_for(seconds(0));
        dur *= 2;
        if (status == future_status::timeout) {
            lock_guard<mutex> lockCout(coutMutex);
            cout << "          " << this_thread::get_id()
                  << " still waiting." << endl;
        }
        if (status == future_status::ready) {
            lock_guard<mutex> lockCout(coutMutex);
            cout << "          " << this_thread::get_id()
                  << " waiting done => Result: " << fut.get() << endl;
        }
    } while (status != future_status::ready);
    const auto end= steady_clock::now();
    lock_guard<mutex> lockCout(coutMutex);
    cout << "          " << this_thread::get_id()
          << " waiting time: " << getDifference(start,end) << " ms" << endl;
}

int main(){

    cout << endl;

    promise<int> prom;
    shared_future<int> future= prom.get_future();
    thread producerThread(producer, move(prom));

    thread consumerThread1(consumer, future, seconds(4));
    thread consumerThread2(consumer, future, seconds(20));
    thread consumerThread3(consumePeriodically, future);
    thread consumerThread4(consumeWithBackoff, future);

    consumerThread1.join();
    consumerThread2.join();
    consumerThread3.join();
    consumerThread4.join();
    producerThread.join();

    cout << endl;

}

```



I create the promise in the main function (line 98), use the promise to create

the associated future (line 99), and move the promise into a separate thread (line 100). I have to move the promise into the thread because it does not support the copy semantic. That is not necessary for the shared futures (lines 102 - 105); they support the copy semantic and can, therefore, be copied.

Before I talk about the work package of the thread, let me say a few words about the auxiliary function `getDifference` (lines 14 - 19). The function takes two-time points and returns the time duration in milliseconds. I will use the function a few times.

What about the five created threads?

- `producerThread` : executes the function `producer` (lines 21 - 25) and publishes its result 2011 after 5 seconds of sleep. This is the result the futures are waiting for.
- `consumerThread1` : executes the function `consumer` (lines 27 - 44). The thread is waiting for 4 seconds at most (line 30) before it continues with its work; This waiting period is not long enough to get the result from the promise.
- `consumerThread2` : executes the function `consumer` (lines 27 - 44). The thread is waiting 20 seconds at most before it continues with its work.
- `consumerThread3` : executes the function `consumePeriodically` (lines 46 - 67). It sleeps for 700 milliseconds (line 50) and asks for the result of the promise (line 60). Because of the `std::chrono::seconds(0)` in line 51, there is no waiting. If the result of the calculation is available, it will be displayed in line 60.
- `consumerThread4` : executes the function `consumeWithBackoff` (lines 69 - 92). It sleeps in the first iteration 1 second and doubles its sleeping period every iteration. Otherwise, its strategy is similar to the strategy of `consumerThread3`.

Now to the synchronization of the program. Both the clock determining the current time and `std::cout` are shared variables, but no synchronization is necessary. Firstly, the method call `std::chrono::steady_clock::now()` is thread-safe (for example in lines 30 and 40); Secondly, the C++ runtime guarantees that the characters will be written *thread-safe* to `std::cout`. I only used a `std::lock_guard` to wrap `std::cout` (for example in lines 22, 27, and 41)

`std::lock_guard` to wrap `std::cout` (for example in lines 32, 37, and 41).

Although the threads write one after the other to `std::cout`, the output is not easy to understand.

The first output is from the promise, with the remaining outputs from the futures. At first `consumerThread4` asks for the result. The output is indented by 8 characters. `consumerThread4` also displays its ID. `consumerThread3` immediately follows. Its output is indented by 4 characters. The output of `consumerThread1` and `consumerThread2` is not indented.

- `consumeThread1`: waits unsuccessfully 4000.18 ms seconds without getting the result.
- `consumeThread2`: gets the result after 5000.3 ms although its waiting duration is up to 20 seconds.
- `consumeThread3`: gets the result after 5601.76 ms. That's about 5600 milliseconds=  $8 * 700$  milliseconds.
- `consumeThread4`: gets the result after 8193.81 ms. To say it differently. It waits 3 seconds too long.