

Loss

Calculate the model's sigmoid cross entropy loss.

Chapter Goals:

- Calculate the model's loss using sigmoid cross entropy

A. Sigmoid cross entropy

The task for our model is to classify input text sequences as either negative (label `0`) or positive (label `1`). This is equivalent to binary classification. As with regular binary classification, we use sigmoid cross entropy to calculate the model's loss.

For an in-depth and intuitive explanation of sigmoid cross entropy and binary classification, check out the [Machine Learning for Software Engineers](#) course on Educative.

Time to Code!

In this chapter you'll be completing the `calculate_loss` function, which calculates model loss based on the outputs of the BiLSTM.

The first step to calculating the model's loss is to first calculate the logits. We can use the `calculate_logits` function we completed in the previous chapter.

Set `logits` equal to `self.calculate_logits` applied with `lstm_outputs`, `batch_size`, and `sequence_lengths` as arguments.

Since we're performing binary classification, we use sigmoid cross entropy for the loss. We also need to convert the integer labels into floats.

Set `float_labels` equal to `tf.cast` applied with `labels` as the first argument and `tf.float32` as the second argument.

Set `batch_loss` equal to `tf.nn.sigmoid_cross_entropy_with_logits` applied with `float_labels` and `logits` for the `labels` and `logits` keyword

arguments, respectively.

The output of the function is the overall aggregate loss. This means we need to sum together each individual sequence's loss in the batch.

Set `overall_loss` equal to `tf.reduce_sum` applied to `batch_loss`. Then return `overall_loss`.

```
import tensorflow as tf
tf_fc = tf.contrib.feature_column

# Text classification model
class ClassificationModel(object):
    # Model initialization
    def __init__(self, vocab_size, max_length, num_lstm_units):
        self.vocab_size = vocab_size
        self.max_length = max_length
        self.num_lstm_units = num_lstm_units
        self.tokenizer = tf.keras.preprocessing.text.Tokenizer(num_words=self.vocab_size)

    def get_gather_indices(self, batch_size, sequence_lengths):
        row_indices = tf.range(batch_size)
        final_indexes = tf.cast(sequence_lengths - 1, tf.int32)
        return tf.transpose([row_indices, final_indexes])

    # Calculate logits based on the outputs of the BiLSTM
    def calculate_logits(self, lstm_outputs, batch_size, sequence_lengths):
        lstm_outputs_fw, lstm_outputs_bw = lstm_outputs
        combined_outputs = tf.concat([lstm_outputs_fw, lstm_outputs_bw], -1)
        gather_indices = self.get_gather_indices(batch_size, sequence_lengths)
        final_outputs = tf.gather_nd(combined_outputs, gather_indices)
        logits = tf.layers.dense(final_outputs, 1)
        return logits

    # Calculate the loss for the BiLSTM
    def calculate_loss(self, lstm_outputs, batch_size, sequence_lengths, labels):
        # CODE HERE
        pass
```

