

# Calculating Loss

Calculate the model's loss based on logits and sparse outputs.

Chapter Goals:

- Calculate the training loss based on the model's logits and final token sequences

## A. Final token sequence

So far, we've used the input sequences and ground truth sequences for training the encoder-decoder model. The final token sequences are used when calculating the loss.

If we view the decoder as a language model, the ground truth sequences act as the language model's input while the final token sequences act as the "correct" output for the language model.

In a language model, we calculate the loss based on the model's logits and the "correct" output. For the encoder-decoder, the loss is calculated in the same way. We use a sparse softmax cross entropy loss based on the logits and final token sequences, and then zero-out the time steps that correspond to padding.

## B. Sequence mask

Previously, we used the `tf.sign` function to create binary sequences that are used to zero-out padding time steps. However, there is another way to create the binary sequences, also known as a *sequence mask*, using just the sequence lengths of the decoder's inputs.

The function that provides this utility is called `tf.sequence_mask`. It takes in a tensor of sequence lengths as its required argument, and returns the corresponding sequence mask.

```
import tensorflow as tf

# Example sequence lengths
```



```

# Example sequence lengths
seq_lens = tf.constant([3, 4, 2])
binary_sequences = tf.sequence_mask(seq_lens)
int_sequences = tf.sequence_mask(seq_lens, dtype=tf.int32)

with tf.Session() as sess:
    binary_array = sess.run(binary_sequences)
    int_array = sess.run(int_sequences)

print(repr(binary_array))
print(repr(int_array))

```



Creating sequence masks from sequence lengths using `tf.sequence_mask`.

Note that the default return type of `tf.sequence_mask` is a boolean tensor. We can set the `dtype` keyword argument to specify a different return type.

## Time to Code!

In this chapter you'll be completing the `calculate_loss` function, which is used to calculate the model's loss.

When calculating the loss, we need to zero-out any time steps that represent padding, i.e. time steps past the actual sequence length. We'll apply `tf.sequence_mask` to the decoder's input sequence lengths ( `dec_seq_lens` ) to obtain the necessary sequence mask.

**Set `binary_sequences` equal to `tf.sequence_mask` applied with `dec_seq_lens` as the required argument. Also set the `dtype` keyword argument to `tf.float32`, so that `binary_sequences` has type `tf.float32`.**

Just like in language modeling, we use sparse softmax cross entropy as the loss function. We also zero-out the loss at time steps that represent padding.

**Set `batch_loss` equal to `tf.nn.sparse_softmax_cross_entropy_with_logits` applied with `decoder_outputs` and `logits` as the `labels` and `logits` keyword arguments, respectively.**

**Set `unpadded_loss` equal to `batch_loss` multiplied by `binary_sequences`.**

The loss that we return will be the per-sequence loss, the average loss for each sequence in the training batch. This is equivalent to taking the entire loss and dividing it by the batch size.

Set `per_seq_loss` equal to `tf.reduce_sum` applied to `unpadded_loss`, divided by `batch_size`. Then return `per_seq_loss`.

```
import tensorflow as tf
tf_fc = tf.contrib.feature_column
tf_s2s = tf.contrib.seq2seq

# Seq2seq model
class Seq2SeqModel(object):
    def __init__(self, vocab_size, num_lstm_layers, num_lstm_units):
        self.vocab_size = vocab_size
        # Extended vocabulary includes start, stop token
        self.extended_vocab_size = vocab_size + 2
        self.num_lstm_layers = num_lstm_layers
        self.num_lstm_units = num_lstm_units
        self.tokenizer = tf.keras.preprocessing.text.Tokenizer(
            num_words=vocab_size)

    # Calculate the model loss
    def calculate_loss(self, logits, dec_seq_lens, decoder_outputs, batch_size):
        # CODE HERE
        pass
```

