

Slices: Termination of Sharing

In this lesson, we will discuss the concept of termination of sharing in slices.

WE'LL COVER THE FOLLOWING

- Slice sharing
- Making a slice longer may terminate sharing
 - Using capacity to determine whether sharing will be terminated
 - Reserving room for elements

Slice sharing

It is possible to access the same elements by more than one slice. For example, the first two of the eight elements below are being accessed through three slices:

```
import std.stdio;

void main() {
    int[] slice = [ 1, 3, 5, 7, 9, 11, 13, 15 ];
    int[] half = slice[0 .. $ / 2];
    int[] quarter = slice[0 .. $ / 4];

    quarter[1] = 0;    // modify through one slice

    writeln(quarter);
    writeln(half);
    writeln(slice);
}
```



Sharing in slices

The effect of the modification to the second element of `quarter` is seen through all slices.

Making a slice longer may terminate sharing

When viewed this way, slices provide shared access to elements. This sharing poses the question of what happens when a new element is added to one of the slices. Since multiple slices can provide access to the same elements, there may not be room to add elements to a slice without stomping on the elements of other slices of the same array.

D disallows element stomping and answers this question by terminating the sharing relationship if there is no room for the new element. The slice that has no room to grow leaves the sharing. When this happens, all the existing elements of that slice are copied to a new place in memory automatically and the slice starts providing access to these new elements.

To see this in action, let's add an element to `quarter` before modifying its second element:

```
quarter ~= 42; // this slice leaves the sharing because there is no room for the new element
quarter[1] = 0; // for that reason this modification does not affect the other slices
```

The output of the program shows that the modification to the `quarter` slice does not affect the others:

```
import std.stdio;

void main() {
    int[] slice = [ 1, 3, 5, 7, 9, 11, 13, 15 ];
    int[] half = slice[0 .. $ / 2];
    int[] quarter = slice[0 .. $ / 4];

    quarter ~= 42; // this slice leaves the sharing because there is no room for the new element
                  // adding 42 to 'quarter' will overwrite the first element of 'half'
    quarter[1] = 0; // for that reason this modification does not affect the other slices

    writeln(quarter);
    writeln(half);
    writeln(slice);
}
```



Explicitly increasing the length of a slice terminates the sharing as well:



```
import std.stdio;

void main() {
    int[] slice = [ 1, 3, 5, 7, 9, 11, 13, 15 ];
    int[] half = slice[0 .. $ / 2];
    int[] quarter = slice[0 .. $ / 4];

    ++quarter.length; // leaves the sharing
    quarter.length += 5; // leaves the sharing

    writeln(quarter);
    writeln(half);
    writeln(slice);
}
```



Increasing the length of a slice

On the other hand, shortening a slice does not affect sharing. Shortening the slice merely means that the slice now provides access to fewer elements:



```
import std.stdio;

void main(){
    int[] a = [ 1, 11, 111 ];
    int[] d = a;

    d = d[1 .. $]; // shortening from the beginning
    d[0] = 42; // modifying the element through the slice

    writeln(a);    // printing the other slice
}
```



Shortening a slice

As can be seen from the output, the modification through `d` is seen through `a`; the sharing is still in effect.

Reducing the length in different ways does not terminate the sharing either:



```
import std.stdio;

void main(){
    int[] a = [ 1, 11, 111 ];
    int[] d = a;
```

```

int[] d = a;
writeln("Before shortening, a :",a);
writeln("Before shortening, d :",d);

d = d[0 .. $ - 1];      // shortening from the end
--d.length;             // same thing
d.length = d.length - 1; // same thing

writeln("After shortening, a :",a);
writeln("After shortening, d :",d);
}

```



Sharing continues even after reducing the length

The sharing of elements is still in effect.

Using capacity to determine whether sharing will be terminated

There are cases when slices continue sharing elements even after an element is added to one of them. This happens when the element is added to the longest slice and there is room at the end of it:

```

import std.stdio;

void main() {
    int[] slice = [ 1, 3, 5, 7, 9, 11, 13, 15 ];
    int[] half = slice[0 .. $ / 2];
    int[] quarter = slice[0 .. $ / 4];

    slice ~= 42;      // adding to the longest slice ...
    slice[1] = 0;     // ... and then modifying an element

    writeln(quarter);
    writeln(half);
    writeln(slice);
}

```



Addition of element to the longest slice

The **capacity** property of slices determines whether the sharing will be terminated if an element is added to a particular slice. **capacity** is actually a function, but this distinction does not have any significance in this discussion. The value of **capacity** has the following meanings:

- When its value is **0**, it means that this is not the *longest original slice*. In this case, adding a *new* element would definitely relocate the elements of the slice and the sharing would terminate.
- When its value is non-zero, it means that this is the *longest original slice*. In this case, **capacity** denotes the *total* number of elements that this slice can hold without the need to be copied. The number of new elements that can be added can be calculated by *subtracting the actual length of the slice from the capacity value* (`slice.capacity - slice.length`). If the length of the slice equals its **capacity**, then it indicates that there is no room for more elements. In this case, the slice will be copied to a new location if one more element is added.

Accordingly, a program that needs to determine whether the sharing will terminate should use logic similar to the following:

```
if (slice.capacity == 0) {
    /* Its elements would be relocated if one
    more element is added to this slice. */
    // ...
} else {
    /* This slice may have room for new elements
    before needing to be relocated. Let's
    calculate how many: */
    auto howManyNewElements = slice.capacity - slice.length;
    // ...
}
```

An interesting corner case is when there is more than one slice comprising all elements. In such a case all slices report to have **capacity**:

```
import std.stdio;

void main() {
    // Three slices to all elements
    int[] s0 = [ 1, 2, 3, 4 ];
    int[] s1 = s0;
    int[] s2 = s0;

    writeln(s0.capacity);
    writeln(s1.capacity);
    writeln(s2.capacity);
}
```



However, as soon as an element is added to one of the slices, the **capacity** of the others drop to **0**:

```
import std.stdio;

void main() {
    // Three slices to all elements
    int[] s0 = [ 1, 2, 3, 4 ];
    int[] s1 = s0;
    int[] s2 = s0;

    s1 ~= 42; // ← s1 becomes the longest

    writeln(s0.capacity);
    writeln(s1.capacity);
    writeln(s2.capacity);
}
```



Change in capacity

Since the slice with the added element is now the longest, it is the only one with capacity greater than 0:

```
0
7      ← now only s1 has capacity
0
```

Reserving room for elements

Both *copying elements* and *allocating new memory* to increase capacity have some cost. For that reason, appending an element can be an expensive operation. When the number of elements to append is known beforehand, it is possible to reserve capacity for the elements using the **.reserve** function:

```
import std.stdio;

void main() {
    int[] slice;

    slice.reserve(24);
    writeln(slice.capacity);

    foreach (element; 0 .. 17) {
```

```
slice ~= element; // these elements will not be moved  
}  
}
```



Reserving room for elements

The elements of slice would be moved only after there are more than 31 elements.

In the next lesson, we will see operations on all elements.