

Data Durability in Replica Sets

In this lesson, we will look at what happens to the data, stored in a database when a secondary node is killed off.

WE'LL COVER THE FOLLOWING ^

- Implementation
 - Creating a replica set
 - Starting the Mongo Shell
 - Using `rs.initiate()`
 - Using `rs.add()`
 - Setting up a Database
 - Quitting the Shell
 - Killing a Secondary Node
 - Checking the Database
- Terminal

Now that you are clear on how replication works and how it is implemented let's look into some other details.

In one of the [previous](#) lessons, what happens when a *primary node* fails.

However, what happens if one of the *secondary* nodes is killed?

We learned [previously](#) that the goal of replication is to ensure that no data is lost if a server fails. Hence, the data is still available even if a secondary node is killed off.

Implementation

Now, let's create a *replica set*, kill a *secondary* node, and then look at its impact on the data stored in the database.

Follow the steps below:

Follow the steps below.

Creating a replica set

As in the [previous](#) lesson, the three `mongod` instances are initiated first.

The three commands, for running the `mongod` instances, have already been run for you in the [terminal](#) below.

To see the details of the commands used to run these `mongod` instances, go to the [following](#) lesson.

Starting the Mongo Shell

Next, connect to one of the running `mongod` processes using the MongoDB shell client.

Call `mongo` to connect to the `mongod` process that is running on port `27017`.

Note: This will act as our *primary* node.

Type the command below on the [terminal](#):

```
mongo
```

Running "mongo" to connect to port 27017

Using `rs.initiate()`

As you did [previously](#), call the following command in the [terminal](#) next:

```
rs.initiate({_id : "rs0",members: [{_id : 0, host : "localhost:27017"}]})
```

Calling rs.initiate() command

At this point, your *replica* set has been initiated.

Using `rs.add()`

As done in the [previous](#) lesson, add the *secondary nodes* into `members` by typing the following commands in the [terminal](#):



```
rs.add("localhost:27018")
rs.add("localhost:27019")
```

Running rs.add() commands for the other two ports

You should see something like this:

```
> rs.initiate({_id : "rs0",members: [{_id : 0, host : "localhost:27017"}]})
{
  "ok" : 1,
  "operationTime" : Timestamp(1564139084, 1),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1564139084, 1),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
rs0:OTHER> rs.add("localhost:27018")
{
  "ok" : 1,
  "operationTime" : Timestamp(1564139090, 1),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1564139090, 1),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
rs0:PRIMARY> rs.add("localhost:27019")
{
  "ok" : 1,
  "operationTime" : Timestamp(1564139091, 1),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1564139091, 1),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
rs0:PRIMARY>
```

Setting up a Database

Next, set up a database with a collection called `users` (refer to the following [lesson](#) to learn how to do that).

Here are the relevant commands, from that [lesson](#), that you can type into the [terminal](#) to set up a database:



```
db
db.createCollection("users");

db.users.insert({"name" : "Nikola Zivkovic", "blog" : "rubikscore.net", "numberOfArticles" :
db.users.find().pretty();
```

You should see the following result:

```
rs0:PRIMARY> use blog
switched to db blog
rs0:PRIMARY> db
blog
rs0:PRIMARY> db.createCollection("users");
{
  "ok" : 1,
  "operationTime" : Timestamp(1564139795, 2),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1564139795, 2),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
rs0:PRIMARY> db.users.insert({"name" : "Nikola Zivkovic", "blog" : "rubikscore.net", "numberOfArticles"
: 10, "company" : "Vega IT"});
WriteResult({ "nInserted" : 1 })
rs0:PRIMARY> db.users.find().pretty();
{
  "_id" : ObjectId("5d3ae113f0dfd53b77b4b636"),
  "name" : "Nikola Zivkovic",
  "blog" : "rubikscore.net",
  "numberOfArticles" : 10,
  "company" : "Vega IT"
}
rs0:PRIMARY>
```

Quitting the Shell

Next, type the following command into the [terminal](#) to quit the `mongo` shell:

```
quit()
```

This command should bring you out of the shell.

```
rs0:PRIMARY> quit()
root@educative:/#
```

Killing a Secondary Node

Once you are out of the shell, type the following command into the [terminal](#) to see the processes that are running and their relevant PIDs:

```
ps
```

Under the `CMD` heading, you will see the three `mongod` processes running; and their respective `PID`s will be displayed under the `PID` heading.

You should see something like this:

```
root@educative:/# ps
  PID TTY          TIME CMD
   29 pts/0        00:00:00 bash
   40 pts/0        00:00:00 bash
   41 pts/0        00:00:00 mongod
   42 pts/0        00:00:00 mongod
   45 pts/0        00:00:00 mongod
  265 pts/0        00:00:00 ps
```

In order to kill a *secondary* node, type the `kill PIDNumber` command, replacing `PIDNumber` with the `PID` of the *second* `mongod` entry, in the output of the `ps` command. Do this in the [terminal](#) below:

```
kill PIDNumber
```

Type the `ps` command into the terminal again; you should be able to see only *two* processes running now.

```
root@educative:/# ps
  PID TTY          TIME CMD
   29 pts/0        00:00:00 bash
   40 pts/0        00:00:00 bash
   41 pts/0        00:00:00 mongod
   42 pts/0        00:00:00 mongod
   45 pts/0        00:00:00 mongod
  265 pts/0        00:00:00 ps
root@educative:/# kill 42
root@educative:/# ps
  PID TTY          TIME CMD
   29 pts/0        00:00:00 bash
   40 pts/0        00:00:00 bash
   41 pts/0        00:00:01 mongod
   45 pts/0        00:00:01 mongod
  267 pts/0        00:00:00 ps
[3]+  Done                  mongod --dbpath /data/db2 --replSet "rs0" --port 27019 > /dev/null
```

Checking the Database

In the last step, run the command `mongo` to connect to the `mongod` process that is running on port `27017` again.

Once you are connected, type the following commands into the [terminal](#):

```
use blog
show collections
db.users.find().pretty();
```



Note: To learn more about these commands, read the following [lesson](#).


You should see that even though one of the MongoDB servers has died, no data has been lost.

```
rs0:PRIMARY> db.users.find().pretty();
{
  "_id" : ObjectId("5d383c360fbea1b82f117d49"),
  "name" : "Nikola Zivkovic",
  "blog" : "rubikscore.net",
  "numberOfArticles" : 10,
  "company" : "Vega IT"
}
```

Terminal

Try out all of the commands mentioned above in the terminal given below!

● Terminal



[Go back](#) (click here to go back above).

In the next lesson, we will look at the concept of *Sharding* in MongoDB.